


Convolution Neural Networks

COS 314

Outline

- History of Convolution Neural Network
- Application Area
- Structure
- Functionality
- Future



Introduction

- Convolutional Neural Networks or CovNets
- Alex Krizhevsky won the ImageNet Challenge – AlexNet
- Image classification
- Classification


Class

Probabilities for each class



Image Classification

- Image converted to a matrix of pixels
- Grayscale images – a single integer 2D matrix with values in the range 0 to 255
- Colour images – 3 stacked (RGB) 2D matrices with values in the range 0 to 255
- Size of the matrix is resolution dependent
- Preprocessing to reduce the size of the matrix







Image Representation



0	0	0	1	1	1	1	0	0	0
0	0	0	1	0	0	1	0	0	0
0	0	0	1	0	0	1	0	0	0
1	1	1	1	1	1	1	1	1	1
1	1	1	1	1	1	1	1	1	1
0	1	0	0	0	0	0	0	1	0
1	1	1	0	0	0	0	1	1	1
1	0	1	0	0	0	0	1	0	1
1	0	1	0	0	0	0	1	0	1
1	0	1	0	0	0	0	1	0	1




Image Representation

Grayscale

8	120	221	189	145	212	33	244
149	20	30	32	205	177	208	173
47	19	3	249	102	229	161	122
0	86	198	197	28	147	107	84
243	40	131	54	150	18	11	158
123	96	137	10	139	111	241	183
239	217	100	153	131	236	230	113
237	166	62	169	129	32	141	185

8 bits per pixel.
Values 0 – 255
Height x Width



Image Representation

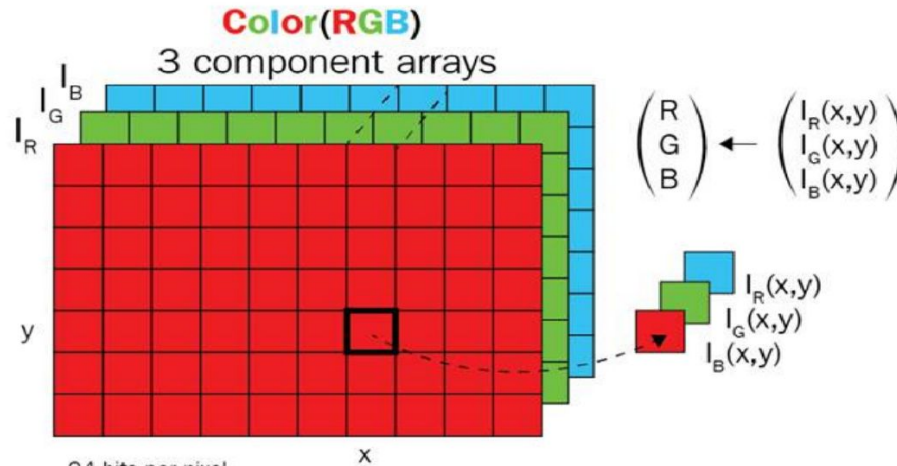


157	153	174	168	160	152	129	161	172	161	155	166
155	182	163	74	75	62	39	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	6	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	225	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

157	153	174	168	160	152	129	161	172	161	155	166
155	182	163	74	75	62	39	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	6	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	225	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218




Image Representation





24 bits per pixel.

RB Red/Green/Blue


Height x Width X 3:



Fully Connected vs Sparse Connected



Parameter Sharing




CNN Layers


- Convolutional layer
- Nonlinear or ReLU layer
- Pooling layer
- Fully connected layer



Overview



Convolution Layer



Convolution Layer

- First layer after the input layer
- Additional convolutional layers
- Convolution operator
- Filter or kernel
- Feature map
- Sparsely connected layer – parameter sharing




Convolution Layer

- Different convolution operators include:
 - Identity
 - Edge detection
- Different convolutions are used with different kernels/filters.
- Selection of values for kernels/filters
- Stride
- Zero padding



Determining the Feature Map

- Decide on the dimensions of the kernel, e.g. 3x3
- Determine the values of the kernel.
- Apply the convolution to produce the feature map.
- Map the kernel across the pixel matrix to produce the feature map.



Determining the Feature Map -2D Example

Input Matrix

0	1	0	1	0	0
0	1	1	0	0	0
1	0	1	1	0	1
1	0	0	1	1	1
0	1	1	1	0	0
0	1	0	1	0	1

6X6

Filter/Kernel

0	0	0
1	0	1
-1	0	-1


3X3

*

=

Output Feature Map

4X4



Determining the Feature Map

Input Matrix

0	1	0	1	0	0
0	1	1	0	0	0
1	0	1	1	0	1
1	0	0	1	1	1
0	1	1	1	0	0
0	1	0	1	0	1

Filter/Kernel

0	0	0
1	0	1
-1	0	-1

*

=

Output Feature Map



Determining the Feature Map

Input Matrix

0	1	0	1	0	0
0	1	1	0	0	0
1	0	1	1	0	1
1	0	0	1	1	1
0	1	1	1	0	0
0	1	0	1	0	1

Filter/Kernel

0	0	0
1	0	1
-1	0	-1

*

=

Output Feature Map

$$\text{Convolution} = 0*0 + 1*0 + 0*0 + 0*1 + 1*0 + 1*1 + 1*-1 + 0*0 + 1*-1 = -1$$



Determining the Feature Map

Input Matrix

0	1	0	1	0	0
0	1	1	0	0	0
1	0	1	1	0	1
1	0	0	1	1	1
0	1	1	1	0	0
0	1	0	1	0	1

Filter/Kernel

0	0	0
1	0	1
-1	0	-1

*

=

Output Feature Map

-1			

$$\text{Convolution} = 1*0 + 0*0 + 1*0 + 1*1 + 1*0 + 0*1 + 0*-1 + 1*0 + 1*-1 = 0$$



Determining the Feature Map

Input Matrix

0	1	0	1	0	0
0	1	1	0	0	0
1	0	1	1	0	1
1	0	0	1	1	1
0	1	1	1	0	0
0	1	0	1	0	1

Filter/Kernel

0	0	0
1	0	1
-1	0	-1

*

=

Output Feature Map

-1	0		

$$\text{Convolution} = 0*0 + 1*0 + 0*0 + 1*1 + 0*0 + 0*1 + 1*-1 + 1*0 + 0*-1 = 0$$



Determining the Feature Map

Input Matrix

0	1	0	1	0	0
0	1	1	0	0	0
1	0	1	1	0	1
1	0	0	1	1	1
0	1	1	1	0	0
0	1	0	1	0	1

Filter/Kernel

0	0	0
1	0	1
-1	0	-1

*

=

Output Feature Map

-1	0	0	

$$\text{Convolution} = 1*0 + 0*0 + 0*0 + 0*1 + 0*0 + 0*1 + 1*-1 + 0*0 + 1*-1 = -2$$



Determining the Feature Map

Input Matrix

0	1	0	1	0	0
0	1	1	0	0	0
1	0	1	1	0	1
1	0	0	1	1	1
0	1	1	1	0	0
0	1	0	1	0	1

Filter/Kernel

0	0	0
1	0	1
-1	0	-1

*

=


Output Feature Map

-1	0	0	-2
-1			


$$\text{Convolution} = 0*0 + 1*0 + 1*0 + 1*1 + 0*0 + 1*1 + 1*-1 + 0*0 + 0*-1 = -1$$



Convolution Layer - 3D Example




ReLU Layer




Nonlinear ReLU Layer

- Is used to incorporate nonlinearity
- The ReLU activation function is used
- $f(x) = \max(0, x)$



Pooling Layer



Pooling Layer

- Reduces the number of parameters
- This is achieved by reducing the size of the feature map

Feature map

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Max pooling




2 x 2 filter

Feature map

6	8
3	4



Fully Connected Layer




Fully Connected Layer

- Output layer
- Last or last two layers
- Softmax activation function
- Transfer learning



Fully Connected Layer



CNN - Architectures

- LeNet
- AlexNet
- GoogLeNet
- VGGNet
- ResNet
- DenseNet
- ImageNet and transfer learning



Image Classification -Transformers

- Use **Attention**
- Natural Language Processing based on RNN
- Vision Transformer performs better than CNNs



Multi-Task Learning

- Multi-task learning (MTL) is a machine learning technique where a model is trained on multiple learning tasks simultaneously.
- The goal is to improve the performance of each task.
- Through leveraging the knowledge and information shared across them.



- Training a single model on multiple related tasks, the model can learn shared representations and features that are beneficial for all tasks.
- This can lead to better performance compared to training separate models for each task.



Advantages of Multitask

- **Improved Generalization:** By considering information from multiple tasks, the model can potentially generalize better to unseen data within each task domain.
- **Data Efficiency:** MTL can be particularly beneficial when dealing with limited data for each individual task. Sharing information across tasks allows the model to learn more effectively with less data.
- **Regularization:** MTL can act as a regularizer, preventing the model from overfitting to the specific features of any single task. (e.g Dropout or early stopping).



QUESTIONS

